


# ML/DL for Everyone Season2

## MNIST Introduction



Code: <https://github.com/deeplearningzerotoall/PyTorch>

Slides: <http://bit.ly/2VrZcWM>

Lecturer: 강현우 znxlwm@gmail.com



# MNIST Introduction

- What is MNIST?
- Code: `mnist_introduction`

# What is MNIST?

## MNIST: handwritten digits dataset



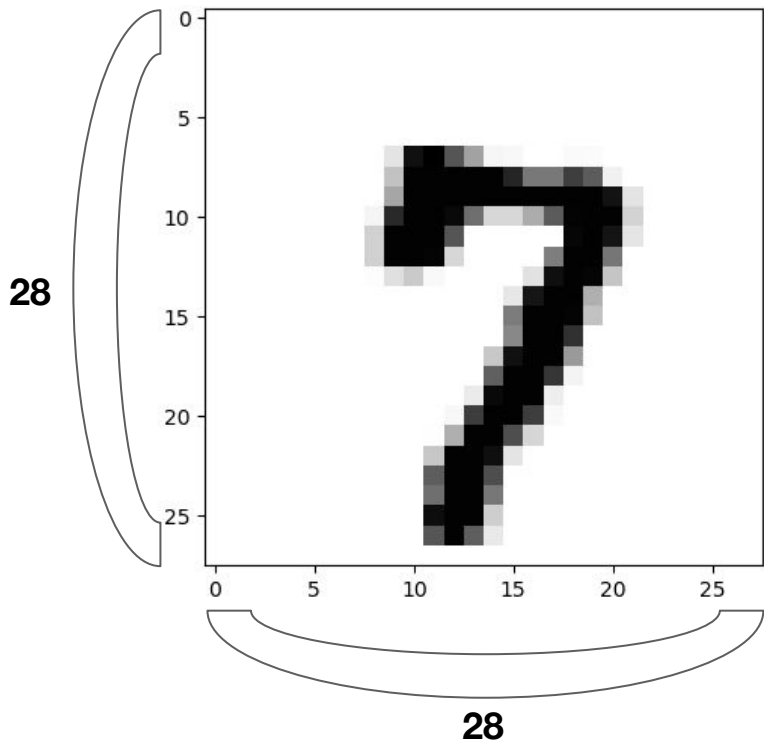
[train-images-idx3-ubyte.gz](#): training set images (9912422 bytes; 60,000 samples)

[train-labels-idx1-ubyte.gz](#): training set labels (28881 bytes)

[T10k-images-idx3-ubyte.gz](#) : test set images (1648877 bytes; 10,000 samples)

[T10k-labels-idx1-ubyte.gz](#) : test set labels (4542 bytes)

# Example of MNIST



- 28 x 28 image
- 1 channel gray image
- 0 ~ 9 digits

```
for X, Y in data_loader:  
    # reshape input image into [batch_size by 784]  
    # label is not one-hot encoded  
    X = X.view(-1, 28 * 28)
```

# torchvision

The torchvision package consists of popular datasets, model architectures, and common image transformations for computer vision.

## torchvision.datasets

- MNIST
- Fashion-MNIST
- EMNIST
- COCO
- LSUN
- ImageFolder
- DatasetFolder
- Imagenet-12
- CIFAR
- STL10
- SVHN
- PhotoTour
- SBU
- Flickr
- VOC

## torchvision.models

- Alexnet
- VGG
- ResNet
- SqueezeNet
- DenseNet
- Inception v3

## torchvision.transforms

- Transforms on PIL Image
- Transforms on  
torch.\*Tensor
- Conversion Transforms
- Generic Transforms
- Functional Transforms

## torchvision.utils

# Reading data

```
import torchvision.datasets as dsets
...
mnist_train = dsets.MNIST(root="MNIST_data/", train=True, transform=transforms.ToTensor(),
download=True)

mnist_test = dsets.MNIST(root="MNIST_data/", train=False, transform=transforms.ToTensor(),
download=True)

data_loader = torch.utils.DataLoader(DataLoader=mnist_train, batch_size=batch_size,
shuffle=True, drop_last=True)
...
for epoch in range(training_epochs):
...
    for X, Y in data_loader:
        # reshape input image into [batch_size by 784]
        # label is not one-hot encoded
        X = X.view(-1, 28 * 28).to(device)
```

# Epoch / Batch size / Iteration

In the neural network terminology:

- one **epoch** = one forward pass and one backward pass of *all* the training examples
- **batch size** = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.
- number of **iterations** = number of passes, each pass using [batch size] number of examples. To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).

**Example:** if you have 1000 training examples, and your batch size is 500, then it will take 2 iterations to complete 1 epoch.

# Softmax

```
# MNIST data image of shape 28 * 28 = 784
linear = torch.nn.Linear(784, 10, bias=True).to(device)
# initialization
torch.nn.init.normal_(linear.weight)
# parameters
training_epochs = 15
batch_size = 100
# define cost/loss & optimizer
criterion = torch.nn.CrossEntropyLoss().to(device) # Softmax is internally defined
optimizer = torch.optim.SGD(linear.parameters(), lr=0.1)

for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = len(data_loader)
    for X, Y in data_loader:
        # reshape input image into [batch_size by 784]
        # Label is not one-hot encoded
        X = X.view(-1, 28 * 28).to(device)
        optimizer.zero_grad()
        hypothesis = linear(X)
        cost = criterion(hypothesis, Y)
        cost.backward()
        avg_cost += cost / total_batch
    print("Epoch: ", "%04d" % (epoch+1), "cost =", "{:.9f}".format(avg_cost))

Epoch: 0001 cost = 2.511683702
Epoch: 0002 cost = 0.977319956
Epoch: 0003 cost = 0.797017217
Epoch: 0004 cost = 0.710427940
Epoch: 0005 cost = 0.655205429
Epoch: 0006 cost = 0.615207732
Epoch: 0007 cost = 0.584421575
Epoch: 0008 cost = 0.559486568
Epoch: 0009 cost = 0.538655698
Epoch: 0010 cost = 0.520880997
Epoch: 0011 cost = 0.505315244
Epoch: 0012 cost = 0.491431117
Epoch: 0013 cost = 0.479477882
Epoch: 0014 cost = 0.468681127
Epoch: 0015 cost = 0.458788306
Learning finished
Accuracy: 0.8718999624252319
```



# Test

*# Test the model using test sets*

**With** torch.no\_grad():

    X\_test = mnist\_test.test\_data.view(-1, 28 \* 28).float().to(device)

    Y\_test = mnist\_test.test\_labels.to(device)

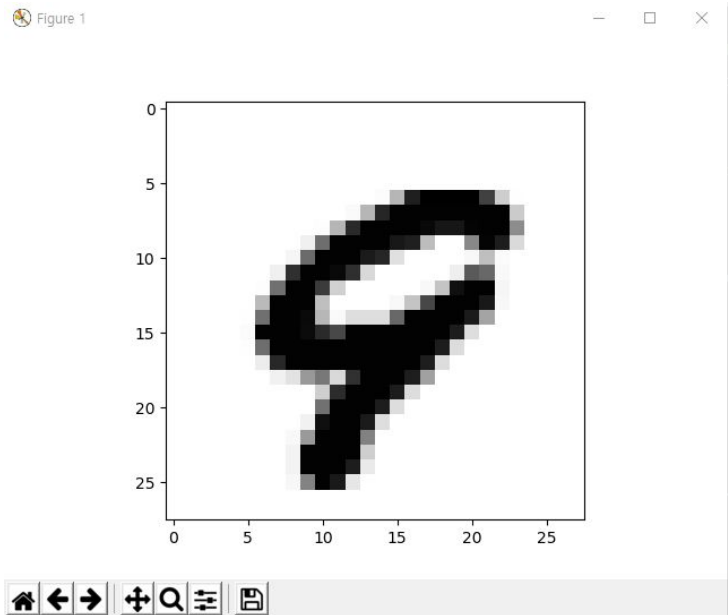
    prediction = linear(X\_test)

    correct\_prediction = torch.argmax(prediction, 1) == Y\_test

    accuracy = correct\_prediction.float().mean()

**print**("Accuracy: ", accuracy.item())

# Visualization



```
import matplotlib.pyplot as plt
import random

...
r = random.randint(0, len(mnist_test) - 1)
X_single_data = mnist_test.test_data[r:r + 1].view(-1, 28 *
28).float().to(device)
Y_single_data = mnist_test.test_labels[r:r + 1].to(device)

print("Label: ", Y_single_data.item())
single_prediction = linear(X_single_data)
print("Prediction: ", torch.argmax(single_prediction,
1).item())

plt.imshow(mnist_test.test_data[r:r + 1].view(28, 28),
cmap="Greys", interpolation="nearest")
plt.show()
```

Label: 8  
Prediction: 8

# What's Next?

- Perceptron